

Shivam More

Lead Data Engineer | Agentic AI Engineer | Cloud & Spark Specialist

Pune, India | Phone: +91-7020093806 | Email: shivamkm1692000@gmail.com | LinkedIn: linkedin.com/in/shivam-more | GitHub: github.com/shivamkm1692000

PROFESSIONAL SUMMARY

Lead Data Engineer with 4+ years of experience building cloud-native data platforms across GCP and Azure for high-volume batch and real-time workloads. Expert in Apache Spark, PySpark, SQL, Airflow, Snowflake, and Databricks, with proven impact in performance, reliability, and cost optimization. Skilled in semantic data modeling, semantic search, and analytics enablement to improve data discoverability and self-service decision-making, and in building production-grade Agentic AI/GenAI systems with LangChain, LangGraph, RAG, and vector databases.

CORE SKILLS

- **Languages & Data:** Python (Pandas, PySpark, FastAPI), SQL (CTEs, window functions, optimization), Spark, Databricks, Delta Lake, Kafka, dbt, Shell, Java
- **Semantic & Analytics:** Semantic data modeling, semantic layer design, semantic search, analytics enablement, data discoverability
- **Cloud & Warehousing:** GCP (BigQuery, Dataproc, GCS, Pub/Sub, Vertex AI), Azure (ADLS Gen2, Synapse, Data Factory, Azure OpenAI, AI Search), Snowflake, PostgreSQL, SQL Server, MongoDB
- **Agentic AI / GenAI:** LangChain, LangGraph, LlamaIndex, RAG, tool calling, MCP, OpenAI, Anthropic Claude, Pinecone, Weaviate, pgvector, FAISS
- **Orchestration & DevOps:** Apache Airflow, Azure Data Factory, Docker, Kubernetes, Jenkins, Azure DevOps, GitHub Actions, MLflow, Azure Monitor, Stackdriver

PROFESSIONAL EXPERIENCE

Persistent Systems, Pune, India

Lead Data Engineer (Lead Software Engineer) • Apr 2026 - Present

- Lead architecture and delivery of cloud-native data and AI platforms on GCP and Azure; set technical roadmap and engineering standards for the team.
- Led semantic data modeling for shared business entities and metrics, improving cross-domain consistency and analytics enablement.
- Architected Agentic AI solutions with LangChain and LangGraph (Azure OpenAI, Vertex AI), integrating RAG, tool calling, and vector search.
- Designed semantic search over structured and unstructured data to accelerate analyst self-service and reduce time-to-insight.

Senior Data Engineer (Senior Software Engineer) • Sep 2024 - Mar 2026

- Reduced cloud compute spend by **20%** across GCP and Azure through lifecycle policies, autoscaling, and spot instances.
- Designed centralized observability with Azure Monitor and Stackdriver, reducing pipeline MTTR by **40%**.
- Re-architected legacy PySpark ETL with partitioning, broadcast joins, and Delta Lake, improving throughput by **3x**.
- Standardized semantic-ready data contracts and reusable transformation models to enable downstream analytics platforms.

Software Engineer (Data Engineer) • Jul 2022 - Aug 2024

- Orchestrated ETL pipelines with Airflow DAGs and Snowflake Tasks, sustaining **99.9%** data availability for daily reporting.
- Built near-real-time pipelines with Spark Structured Streaming and Pub/Sub, reducing insight latency from hours to minutes.
- Modeled domain datasets for analytics enablement, improving BI metric trust and reducing ambiguity.

Software Engineering Intern • Feb 2022 - Jul 2022

- Trained on Azure, Databricks, and PySpark; supported migration of on-premise workflows to cloud-native architectures.

SELECTED PROJECTS

- **Agentic RAG Assistant:** LangGraph agent combining semantic search over pgvector with BigQuery SQL tools for analyst self-service Q&A and analytics enablement.
- **Self-Healing Pipeline Agent:** Airflow monitoring agent that classifies failures via LLM tool calls and auto-creates Jira tickets with proposed fixes.

CERTIFICATIONS

- Google Cloud Professional Data Engineer | Google Cloud Associate Cloud Engineer | Microsoft Azure Fundamentals (AZ-900)

EDUCATION

B.E. in Computer Engineering, Savitribai Phule Pune University | 2019 - 2022 | CGPA: 9.39